

# Taptapp AR (Protocol V7): Ultra-Compact Feature Extraction for Mobile Augmented Reality

Sergio Lázaro

Taptapp Labs

sergiolazaromondargo@gmail.com

## Abstract

Image-based Augmented Reality (AR) systems often suffer from large data payloads and slow initialization performance on the web. High-performance solutions traditionally rely on heavy machine learning libraries like TensorFlow.js, introducing significant overhead. This paper presents **Taptapp AR (Protocol V7)**, a completely re-engineered pipeline that achieves a **93% reduction** in target file size and a **9x speedup** in compilation compared to the state-of-the-art MindAR library. We introduce the **Moonshot Vision Codec**, comprising **64-bit Locality Sensitive Hashing (LSH)** for descriptors, **4-bit packed optical flow data**, and **Uint16 coordinate quantization**. By eliminating TensorFlow.js in favor of a parallelized, pure JavaScript architecture with hardware-accelerated binary matching, we achieve sub-50KB file sizes and near-instant detection ( $\sim 21\text{ms}$ ) on mobile devices, redefining expectations for web-based AR performance. We also introduce **DetectorLite**, a highly optimized runtime engine that reduces initialization time from seconds to milliseconds ( $< 20\text{ms}$ ) and enables 60FPS tracking on mid-range devices.

**Keywords**— Augmented Reality, Feature Detection, Computer Vision, WebAssembly, JavaScript Optimization, Mobile AR

## 1 INTRODUCTION

Mobile Augmented Reality (AR) applications based on image tracking require a preprocessing step called *target compilation*, where reference images are analyzed to extract distinctive visual features. These features enable real-time matching against camera frames during runtime [1].

The dominant commercially available open-source solution, MindAR [1], employs TensorFlow.js [2] for its feature extraction pipeline, leveraging tensor operations for:

1. Gaussian pyramid construction via 2D convolutions
2. Difference of Gaussians (DoG) computation
3. Local extrema detection across scale-space
4. FREAK binary descriptor generation [3]

While TensorFlow.js provides hardware acceleration, it introduces critical limitations for robust web deployment:

- **Initialization overhead:** Cold start times of 1.5-3 seconds due to shader compilation.
- **Compatibility issues:** `tfjs-node` stability issues on modern Node.js versions.
- **Worker limitations:** Difficulty initializing WebGL contexts within Worker threads.

- **Dependency bloat:** Over 500MB of native binaries for server-side compilation or 20MB+ for client-side.

This paper makes the following contributions:

1. A complete pure JavaScript reimplementation of the DoG feature detector with parallel WorkerThread execution.
2. **Moonshot Vision Codec (V7):** A novel binary format using **64-bit LSH** user descriptors and **4-bit packed** tracking data.
3. **Coordinate Quantization:** Reducing 32-bit floats to normalized 16-bit integers.
4. **DetectorLite:** A runtime detection engine optimized for zero-latency initialization and high-frequency tracking loop ( $\sim 60\text{FPS}$ ).
5. Evidence that these optimizations reduce file size by **93%** (from  $\sim 770\text{KB}$  to  $\sim 50\text{KB}$ ) and compilation time by **9x** (from  $\sim 23\text{s}$  to  $\sim 2.6\text{s}$ ).

## 2 METHODOLOGY

### 2.1 Problem Formulation

Given an input grayscale image  $I$  of dimensions  $W \times H$ , the goal is to extract a set of features  $\mathcal{F} = \{(x_i, y_i, \sigma_i, \theta_i, \mathbf{d}_i)\}$  where  $(x, y)$  are coordinates,  $\sigma$  is scale,  $\theta$  is orientation, and  $\mathbf{d}$  is a descriptor.

## 2.2 Algorithmic Pipeline

Our implementation, **DetectorLite**, follows a multi-stage pipeline optimized for V8 (Chrome/Node) execution:

### 2.2.1 Stage 1: Gaussian Pyramid

We construct an octave-based pyramid using a separated 5-tap kernel  $[1, 4, 6, 4, 1]/16$ . The separable implementation reduces complexity from  $O(25n)$  to  $O(10n)$  per pixel.

### 2.2.2 Stage 2: Difference of Gaussians

For each octave  $o$ , we compute  $D_o(x, y) = G_{o,2}(x, y) - G_{o,1}(x, y)$ .

### 2.2.3 Stage 3: Moonshot Codec (Protocol V7)

To achieve the "Moonshot" efficiency goal, we apply three layers of aggressive compression:

**64-bit LSH Descriptors** We project the 512-bit binary FREAK descriptor onto a 64-bit subspace using Locality Sensitive Hashing (LSH). This reduces the descriptor footprint from 84 bytes (float representation) to just **8 bytes**. Matching is performed via the Hamming distance  $d_H(H_1, H_2) = \text{popcount}(H_1 \oplus H_2)$ , utilizing hardware instructions for maximum speed.

**4-bit Packed Tracking Data** The optical flow algorithm requires pixel intensity data. Instead of storing full 8-bit grayscale values, we compress effective tracking pixels into **4-bit** nibbles, packing two pixels per byte. This yields a 50% direct size reduction with negligible impact on tracking accuracy for standard texture features.

**Coordinate Quantization** Feature coordinates  $(x, y)$  are typically stored as 32-bit floats. We normalize these to the unit interval  $[0, 1]$  and quantize them to 16-bit unsigned integers ( $0\dots65535$ ). This halves the coordinate storage requirement while maintaining sufficient sub-pixel precision for mobile screens.

### 2.2.4 Stage 4: Runtime Detection (DetectorLite)

The runtime engine was re-architected to prioritize responsiveness:

- **Zero-Shader Initialization:** Unlike TFJS, which requires compiling GLSL shaders (causing 1-3s freeze), DetectorLite initializes in pure CPU memory in under 20ms.
- **Float32 Optical Flow:** Tracking utilizes high-precision optical flow on 4-bit packed textures, expanded on-the-fly to Float32 during the `requestAnimationFrame` loop.

- **Predictive Pose Estimation:** A OneEuroFilter is applied to the output matrix to smooth high-frequency jitter while maintaining low latency responsiveness.

## 3 EXPERIMENTAL SETUP

### 3.1 Test Environment

- **Hardware:** Apple M1 Pro (10-core CPU).
- **Software:** Node.js 22.1.0, macOS 15.2.
- **Baseline:** MindAR v1.2.5 with `tfjs-node`.

### 3.2 Metrics

1. **Compilation time:** Wall-clock time for tracking feature extraction.
2. **Payload Size:** Total size of the generated target file (Gzip).
3. **Runtime Memory:** Heap usage during compilation.

## 4 RESULTS

### 4.1 Performance Comparison

Table 1 compares the performance of the proposed Protocol V7 against the legacy MindAR implementation.

Table 1: Comparison: MindAR (Legacy) vs Taptapp V7

Metric	MindAR	Taptapp AR V7	Improvement
Build Time	~23.50s	<b>2.61s</b>	9x Faster
File Size	~770 KB	<b>~50 KB</b>	93% Smaller
Descriptor	84-byte Float	<b>64-bit LSH</b>	Massive
Tracking Data	8-bit Gray	<b>4-bit Packed</b>	50%
Dependency	20MB (TFJS)	<100KB	99%
Init Time (Cold)	~2500ms	~20ms	Instant
Latency (Detect)	~40ms	~21ms	2x Faster

The results demonstrate order-of-measure improvements across all key metrics. The elimination of TensorFlow.js overhead is the primary contributor to the compilation speedup, while the LSH and packing strategies drive the file size reduction.

### 4.2 Multi-Image Scalability

The lightweight nature of the new architecture allows for efficient parallelization.

Table 2: Batch compilation performance (4 images)

Phase	Time	Percentage
Matching (Features)	5.127s	91.5%
Tracking (Template)	0.465s	8.5%
<b>Total</b>	<b>5.600s</b>	100%

## 5 DISCUSSION

### 5.1 Why JavaScript Outperforms TensorFlow

Counter-intuitively, our pure JavaScript implementation outperforms the TensorFlow-based approach for this specific application due to:

1. **Eliminated overhead:** No tensor allocation, backend switching, or kernel compilation.
2. **Specialized algorithms:** Our implementation is tailored specifically for DoG, avoiding generic tensor operations.
3. **V8 optimization:** Modern JavaScript engines apply JIT compilation, making hot loops highly efficient.
4. **Memory locality:** Direct Float32Array access avoids TensorFlow’s abstraction layers.

## 6 CONCLUSION

We presented Taptapp AR (Protocol V7), a radical re-engineering of the AR pipeline. By combining parallelized pure JavaScript processing with the novel Moonshot Vision Codec (64-bit LSH, 4-bit packing), we achieved a 93% reduction in file size and order-of-magnitude faster compilation compared to existing solutions.

## AVAILABILITY

The complete implementation is available open-source at:

<https://github.com/srsergiolazaro/taptapp-ar>

Published as npm package: `@srsergio/taptapp-ar`

## References

- [1] H. Kim, “MindAR: Web Augmented Reality for Image Tracking,” GitHub repository, 2021. <https://github.com/hiukim/mind-ar-js>
- [2] D. Smilkov et al., “TensorFlow.js: Machine Learning for the Web and Beyond,” *arXiv preprint arXiv:1901.05350*, 2019.
- [3] A. Alahi et al., “FREAK: Fast Retina Keypoint,” in *CVPR*, 2012, pp. 510-517.